

# Intro to deep learning




Yandex  
Data Factory

LAMBDA 




**British Hedgehog  
Preservation Society**


# Recap: logistic regression



# Nonlinear dependencies



What we have



What we want


- How to get that?

# Feature extraction

Loss, for example:

$$L(y, y_{pred}) = y \cdot \log y_{pred} + (1 - y) \cdot \log (1 - y_{pred})$$

Model:



Training:


$$\operatorname{argmin}_{\theta} L(y, y_{pred}(X, \theta))$$

# Feature extraction

Loss, for example:


$$L(y, y_{pred}) = y \cdot \log y_{pred} + (1 - y) \cdot \log (1 - y_{pred})$$

Model:



Gradient:

$$\frac{\delta L(y, y_{pred}(X, \theta_1))}{\delta \theta_1}$$




Features would tune to your problem automatically!

# What do we want, exactly?

Loss, for example:

$$L(y, y_{pred}) = y \cdot \log y_{pred} + (1 - y) \cdot \log (1 - y_{pred})$$

Model:



Training:

?


$$\operatorname{argmin}_{\theta_1} L(y, y_{pred}(X, \theta_1, \theta_2))$$

# What do we want, exactly?

Loss, for example:

$$L(y, y_{pred}) = y \cdot \log y_{pred} + (1 - y) \cdot \log (1 - y_{pred})$$

Model:




Gradients:

$$\frac{\partial L(y, y_{pred}(X, \theta_1, \theta_2))}{\partial \theta_2} \quad \frac{\partial L(y, y_{pred}(X, \theta_1, \theta_2))}{\partial \theta_1}$$



# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$y_{pred} = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Try linear


$$y_{pred} = \sigma \left( \sum_j w_j^o \left( \sum_i w_{ij}^h x_i + b_j^h \right) + b^o \right)$$

$$w'_i = \sum_j w_j^o w_{ij}^h \qquad b' = \sum_j w_j^o b_j^h + b^o$$

$$y_{pred} = \sigma \left( \sum_i w'_i x_i + b' \right)$$

# Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$


Output:

$$y_{pred} = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

# Nonlinearity

Model:




$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right) \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Gradients:


$$\frac{\delta L(y, y_{pred}(X, w_j^o, b^o, w_{ij}^h, b_j^h))}{\delta w_j^o, \delta b^o, \delta w_{ij}^h, \delta b_j^h}$$

# Nonlinearity

- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$
- $f(a) = \max(0, a)$
- $f(a) = \log(1+e^x)$




# Biological inspiration






# Biological inspiration



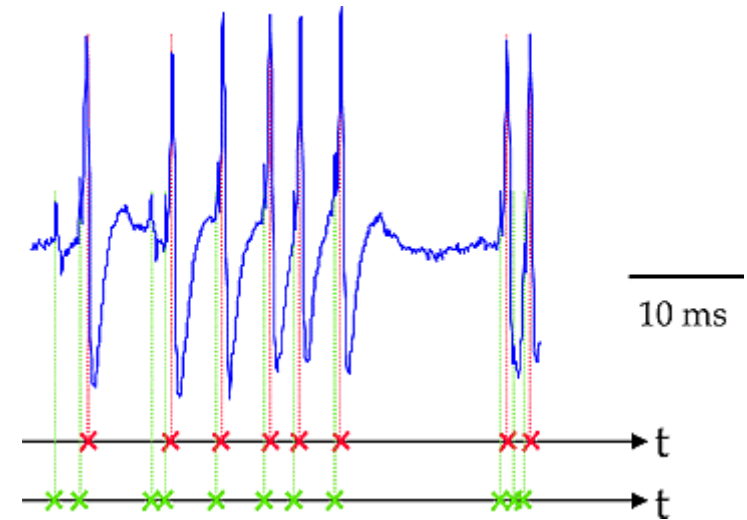
# Biological inspiration



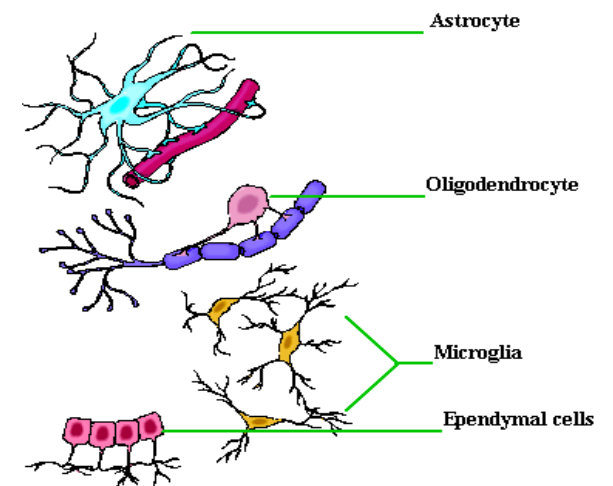


# Not actual neurons :)

- Neurons react in “spikes”, not real numbers
- Neurons maintain/change their states over time
- No one knows for sure how they “train”
- Neuroglial cells are important  
But noone knows, why




Neuroglial Cells of the CNS



# Connectionist phrasebook


- Layer – a building block for NNs :
  - “Dense layer”:  $f(x) = Wx + b$
  - “Nonlinearity layer”:  $f(x) = \sigma(x)$
  - Input layer, output layer
  - A few more we gonna cover later
- Activation – layer output
  - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

# Connectionist phrasebook



- “Train it via backprop!”

# Connectionist phrasebook




How do we train it?




**Discrete Choices**


⋮



**Layer 2 Features**



**Layer 1 Features**



**Original Data**


# Potential caveats?

# Potential caveats?


- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

# Regularization

- L1, L2, as usual
- Dropout



(a) Standard Neural Net



(b) After applying dropout.



# Computation




*The Difference between a CPU and GPU*



**CPU**

MULTIPLE CORES



**GPU**

THOUSAND OF CORES

# Application: Image recognition



“Dog”

# Application: Image recognition




“Dog”

<a particular kind  
of dog>

“Dog tongue”


“Animal sadism”

# Application: Image recognition




$$K_S^0 \rightarrow \pi^+ \pi^-$$


# Classical approach



# NN approach




# Convolutional NNs



Intuition: how cat-like is this square?




# Convolutional NNs



Intuition: how cat-like is this square?

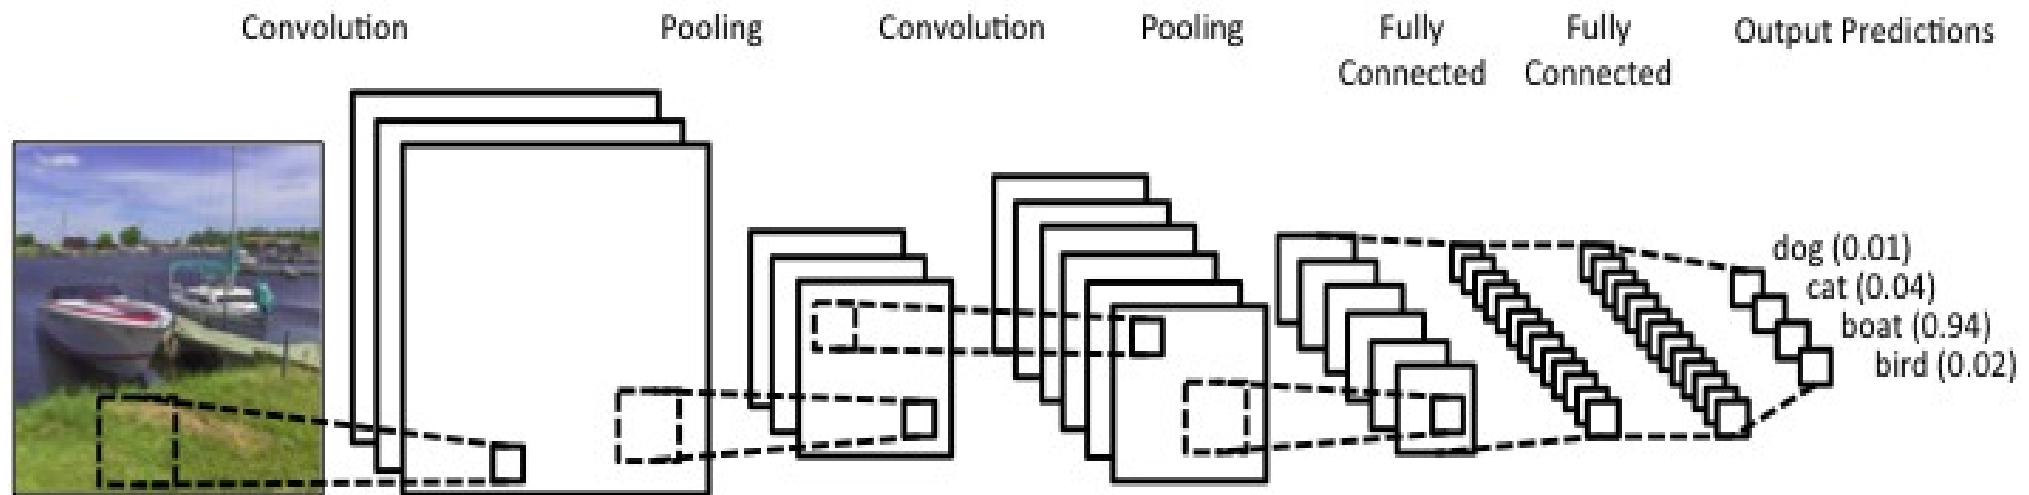


# Convolutional NNs



Intuition: What is the max cat-likelihood over this area?


# Convolutional NNs



# Gradient descent

## Gradient descent algorithm

repeat until convergence {  
     $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$   
    (for  $j = 1$  and  $j = 0$ )  
}




# Newton-Raphson

Parameter update

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma [\mathbf{H}f(\mathbf{x}_n)]^{-1} \nabla f(\mathbf{x}_n).$$

Hessian:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$



Red: Newton-Raphson  
Green: gradient descent

Any drawbacks?

# Newton-Raphson

Parameter update


$$\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma [\mathbf{H}f(\mathbf{x}_n)]^{-1} \nabla f(\mathbf{x}_n).$$

Hessian:

$$\begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$



$$\begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$



Red: Newton-Raphson  
Green: gradient descent

Impractical for large NNs

# SGD with momentum

Idea: move towards “overall gradient direction”,  
Not just current gradient.

$$\Delta w := \eta \nabla Q_i(w) + \alpha \Delta w$$

$$w := w - \Delta w$$

# AdaGrad

Idea: decrease learning rate individually for each parameter in proportion to sum of it's gradients so far.

Let  $g_{\tau,j} = \frac{\delta L}{\delta w_j}$  on  $\tau_{th}$  tick

$$G_{j,j} = \sum_{\tau=1}^t g_{\tau,j}^2$$

$$w_j := w_j - \frac{\eta}{\sqrt{G_{j,j}}} g_j.$$

# RMSProp


Idea: make sure all gradient steps have approximately same magnitude (by keeping moving average of magnitude)

$$v(w, t) := \gamma v(w, t - 1) + (1 - \gamma)(\nabla Q_i(w))^2$$

$$w := w - \frac{\eta}{\sqrt{v(w, t)}} \nabla Q_i(w)$$



# Alltogether



# Moar stuff

- Adadelta
  - Adam
  - Adamax
- Hessian-free
- Nesterov-momentum

# Nuff

Let's code some neural networks!